

I got the following error on ubuntu when using pytorch, can you help : "" File
"/home/gongai/.local/lib/python3.10/site-packages/torch/cuda/init.py", line 298, in
lazyinit torch._C._cuda_init() RuntimeError: CUDA unknown error - this may be due to an
incorrectly set up environment, ""

Claude

1) check if your NVIDIA drivers are properly installed:

``

nvidia-smi

``

2) Verify your CUDA version:

``

\$ nvcc --version

nvcc: NVIDIA (R) Cuda compiler driver

Copyright (c) 2005-2023 NVIDIA Corporation

Built on Fri_Nov__3_17:16:49_PDT_2023

Cuda compilation tools, release 12.3, V12.3.103

Build cuda_12.3.r12.3/compiler.33492891_0

``

3) Check which version of PyTorch you have installed:

``

\$ python3 -c "import torch; print(torch.__version__)"

2.1.2+cu121

``

```
(promptfix) gongai@ducklover1:~$ nvidia-smi
```

```
Sat Oct 26 10:52:15 2024
```

NVIDIA-SMI 550.107.02				Driver Version: 550.107.02				CUDA Version: 12.4			
GPU	Name		Persistence-M	Bus-Id	Disp.A	Volatile	Uncorr. ECC				
Fan	Temp	Perf	Pwr:Usage/Cap		Memory-Usage	GPU-Util	Compute M.				
							MIG M.				
0	NVIDIA	GeForce RTX 4060	Off	00000000:01:00.0	On		N/A				
0%	36C	P8	N/A / 115W	366MiB /	8188MiB	0%	Default				N/A

Processes:							
GPU	GI	CI	PID	Type	Process name	GPU Memory Usage	
	ID	ID					
0	N/A	N/A	9950	G	/usr/lib/xorg/Xorg	116MiB	
0	N/A	N/A	10087	G	/usr/bin/gnome-shell	140MiB	
0	N/A	N/A	12191	G	...seed-version=20241025-130107.645000	98MiB	

```
(promptfix) gongai@ducklover1:~$ nvcc --version
```

```
nvcc: NVIDIA (R) Cuda compiler driver
```

```
Copyright (c) 2005-2023 NVIDIA Corporation
```

```
Built on Fri Nov 3 17:16:49 PDT 2023
```

```
Cuda compilation tools, release 12.3, V12.3.103
```

```
Build cuda_12.3.r12.3/compiler.33492891_0
```

```
(promptfix) gongai@ducklover1:~$ python3 -c "import torch; print(torch.__version__)"
```

```
2.1.2+cu121
```

<https://pytorch.org/>

NOTE: Latest PyTorch requires Python 3.9 or later.

PyTorch Build	Stable (2.5.0)			Preview (Nightly)	
Your OS	Linux		Mac	Windows	
Package	Conda	Pip		LibTorch	Source
Language	Python			C++ / Java	
Compute Platform	CUDA 11.8	CUDA 12.1	CUDA 12.4	ROCm 6.2	CPU
Run this Command:	pip3 install torch torchvision torchaudio				

```
^^^
```

```
pip3 install torch torchvision torchaudio
```

```
^^^
```

Instead of installing the version (which is old) you suggested, I went to pytorch.org, and ran this command to install the latest pytorch 2.5: `pip3 install torch torchvision torchaudio`, then verify its version by `python3 -c "import torch; print(torch.__version__); print(torch.cuda.is_available())"`, it throws the following error: `Traceback (most recent call last): File "<string>", line 1, in <module> File "/home/gongai/anaconda3/envs/promptfix/lib/python3.10/site-packages/torch/init.py", line 235, in <module> from torch._C import * # noqa: F403 ImportError: libcudnn.so.8: cannot open shared object file: No such file or directory`, libcudnn.so.8 share lib is missing, how to install that?

Need to install cuDNN (CUDA Deep Neural Network library)

...

```
sudo apt update
```

```
sudo apt install nvidia-cuda-toolkit
```

```
# Add NVIDIA repository key
```

```
wget https://developer.download.nvidia.com/compute/cuda/repos/ubuntu2204/x86_64/cuda-ubuntu2204.pin
```

```
sudo mv cuda-ubuntu2204.pin /etc/apt/preferences.d/cuda-repository-pin-600
```

```
sudo apt-key adv --fetch-keys
```

```
https://developer.download.nvidia.com/compute/cuda/repos/ubuntu2204/x86_64/3bf863cc.pub
```

```
sudo add-apt-repository "deb
```

```
https://developer.download.nvidia.com/compute/cuda/repos/ubuntu2204/x86_64/ "
```

```
# Install cuDNN
```

```
sudo apt update
```

```
sudo apt install libcudnn8
```

```
python3 -c "import torch; print(torch.__version__); print(torch.cuda.is_available())"
```

...

New error:

Traceback (most recent call last):

File "<string>", line 1, in <module>

File "/home/gongai/anaconda3/envs/promptfix/lib/python3.10/site-packages/torch/_init_.py", line 235, in <module>

from torch._C import * # noqa: F403

ImportError: libnccl.so.2: cannot open shared object file: No such file or directory

You need to install the NVIDIA Collective Communications Library (NCCL).

...

```
# Install NCCL library
```

```
sudo apt update
```

```
sudo apt install libnccl2 libnccl-dev
```

```
$ python3 -c "import torch; print(torch.__version__); print(torch.cuda.is_available())"
```

```
2.1.2+cu121
```

```
True
```

...

running `scripts/inference.py` throws this error:

```
File "/home/gongai/anaconda3/envs/promptfix/lib/python3.10/site-packages/torch/cuda/amp/autocast_mode.py", line 119, in
decorate_fwd
    return fwd(*_cast(args, cast_inputs), **_cast(kwargs, cast_inputs))
File "/home/gongai/projects/wgong/PromptFix/./stable_diffusion/ldm/modules/diffusionmodules/util.py", line 145, in forward
    output_tensors = ctx.run_function(*ctx.input_tensors)
File "/home/gongai/projects/wgong/PromptFix/./stable_diffusion/ldm/modules/attention.py", line 251, in _forward
    x = self.attn1(self.norm1(x)) + x
File "/home/gongai/anaconda3/envs/promptfix/lib/python3.10/site-packages/torch/nn/modules/module.py", line 1518, in wrapped_call_impl
    return self._call_impl(*args, **kwargs)
File "/home/gongai/anaconda3/envs/promptfix/lib/python3.10/site-packages/torch/nn/modules/module.py", line 1527, in _call_impl
    return forward_call(*args, **kwargs)
File "/home/gongai/projects/wgong/PromptFix/./stable_diffusion/ldm/modules/attention.py", line 192, in forward
    sim = einsum('b i d, b j d -> b i j', q, k) * self.scale
File "/home/gongai/anaconda3/envs/promptfix/lib/python3.10/site-packages/torch/functional.py", line 377, in einsum
    return VF.einsum(equation, operands) # type: ignore[attr-defined]
torch.cuda.OutOfMemoryError: CUDA out of memory. Tried to allocate 6.19 GiB. GPU 0 has a total capacity of 7.74 GiB of which 1.64 GiB is free. Including non-PyTorch memory, this process has 5.58 GiB memory in use. Of the allocated memory 5.03 GiB is allocated by PyTorch, and 430.07 MiB is reserved by PyTorch but unallocated. If reserved but unallocated memory is large try setting max_split_size_mb to avoid fragmentation.  See documentation for Memory Management and PYTORCH_CUDA_ALLOC_CONF
```